

---

# **VNIT Housing Price Prediction Documentation**

***Release 1.0.0***

**Shivam Awasthi, Abhishek Bhadauria**

**Dec 07, 2025**



## **Contents**



Welcome to the VNIT Housing Price Prediction project documentation. This comprehensive guide covers the complete machine learning pipeline including data preprocessing, exploratory data analysis, model optimization, and deployment.



# 1

## Project Overview

The VNIT Housing Price Prediction project is an end-to-end machine learning solution that predicts house prices using advanced regression techniques. The project implements best practices in data science including:

- **Data Preprocessing & EDA** - Comprehensive data cleaning and exploration
- **Hyperparameter Optimization** - Optuna-based parameter tuning
- **Model Training** - XGBoost with cross-validation
- **Experiment Tracking** - MLflow for reproducibility
- **Performance Monitoring** - Evidently for model health checks
- **Workflow Orchestration** - Prefect for pipeline management



# 2

## Key Metrics

- **Model R<sup>2</sup> Score:** 0.8990
- **RMSE:** \$27,836.04
- **MAE:** \$17,279.64
- **Optimization Trials:** 50 (Optuna)
- **Cross-Validation Folds:** 5



## Quick Start

### 1. Install Dependencies:

```
pip install -r requirements.txt
```

### 2. Run Preprocessing Pipeline:

```
=====
VNIT Housing Price Prediction
=====

.. image:: https://img.shields.io/badge/Python-3.8+-blue.svg
.. image:: https://img.shields.io/badge/License-MIT-green.svg

Welcome to the VNIT Housing Price Prediction project documentation. This comprehensive guide covers the complete machine learning pipeline including data preprocessing, exploratory data analysis, model optimization, and deployment.
```

#### Project Overview

```
=====
```

The VNIT Housing Price Prediction project is an end-to-end machine learning solution that predicts house prices using advanced regression techniques. The project implements best practices in data science including:

- \*\*Data Preprocessing & EDA\*\* - Comprehensive data cleaning and exploration
- \*\*Hyperparameter Optimization\*\* - Optuna-based parameter tuning
- \*\*Model Training\*\* - XGBoost with cross-validation

(continues on next page)

(continued from previous page)

- **Experiment Tracking** - MLflow for reproducibility
- **Performance Monitoring** - Evidently for model health checks
- **Workflow Orchestration** - Prefect for pipeline management

## Key Metrics

---

- **Model R<sup>2</sup> Score**: 0.8990
- **RMSE**: \$27,836.04
- **MAE**: \$17,279.64
- **Optimization Trials**: 50 (Optuna)
- **Cross-Validation Folds**: 5

## Quick Start

---

### 1. **Install Dependencies**:

```
pip install -r requirements.txt
```

### 2. **Run Preprocessing Pipeline**:

```
cd src/preprocessing  
python preprocessing.py ../../data/train.csv
```

### 3. **Train Model**:

```
cd src/modeling  
python modeling.py
```

### 4. **Launch Streamlit App**:

```
streamlit run src/app/streamlit_app.py
```

## Project Structure

---

::

```
VNIT_project/  
  └── data/                                # Dataset files  
      ├── train.csv                          # Training data  
      └── test.csv                           # Test data
```

(continues on next page)

(continued from previous page)

X_train.csv	# Preprocessed features
data_description.txt	# Feature descriptions
models/	# Trained models
└── xgboost_model.joblib	# Serialized XGBoost model
src/	
└── preprocessing/	# Data preprocessing modules
├── preprocessing.py	# Preprocessing pipeline
├── monitoring.py	# Data quality monitoring
└── EDADataPreProcessing.ipynb	
└── modeling/	# Model training modules
├── modeling.py	# Model pipeline
└── Modeling.ipynb	
└── app/	# Streamlit application
└── streamlit_app.py	
└── mlruns/	# MLflow experiment tracking
docs/	# Documentation
Dockerfile	# Docker configuration
docker-compose.yml	# Docker Compose setup
requirements.txt	# Python dependencies

## Table of Contents

=====

```
.. toctree::
:maxdepth: 2
:caption: Getting Started

getting_started
installation
quick_start

.. toctree::
:maxdepth: 2
:caption: Modules & APIs

modules/preprocessing
modules/monitoring
modules/modeling

.. toctree::
:maxdepth: 2
:caption: Guides
```

(continues on next page)

(continued from previous page)

```
guides/data_preprocessing  
guides/model_training  
guides/deployment  
  
.. toctree:::  
    :maxdepth: 2  
    :caption: Reference  
  
    api_reference  
    troubleshooting  
    faq
```

## Features

---

- ✓ \*\*Automated Data Pipeline\*\* - End-to-end preprocessing with Prefect
- ✓ \*\*Hyperparameter Optimization\*\* - Optuna integration for parameter tuning
- ✓ \*\*Experiment Tracking\*\* - MLflow for model versioning and comparison
- ✓ \*\*Data Quality Monitoring\*\* - Evidently for drift detection
- ✓ \*\*Model Evaluation\*\* - Comprehensive metrics and validation
- ✓ \*\*Production Ready\*\* - Docker support and API deployment
- ✓ \*\*Documentation\*\* - Complete Sphinx documentation with examples

## Technologies Used

---

- \*\*ML Framework\*\*: XGBoost
- \*\*Hyperparameter Tuning\*\*: Optuna
- \*\*Experiment Tracking\*\*: MLflow
- \*\*Workflow Orchestration\*\*: Prefect
- \*\*Data Quality\*\*: Evidently
- \*\*Data Processing\*\*: Pandas, NumPy, Scikit-learn
- \*\*Visualization\*\*: Matplotlib, Seaborn
- \*\*Deployment\*\*: Streamlit, Docker

## Performance

---

The final XGBoost model achieves excellent performance on the test set:

(continues on next page)

(continued from previous page)

- **\*\*R<sup>2</sup> Score\*\*:** 0.8990 (explains 89.90% of variance)
- **\*\*Root Mean Squared Error\*\*:** \$27,836
- **\*\*Mean Absolute Error\*\*:** \$17,280
- **\*\*Training Samples\*\*:** 1,168
- **\*\*Test Samples\*\*:** 292

## Contributing

---

---

Contributions are welcome! Please feel free to submit pull requests ↗ or open issues for bugs and feature requests.

## Authors

---

---

- **\*\*Shivam Awasthi\*\*** (Data Science Engineer)
- **\*\*Abhishiek Bhadauria\*\*** (Data Science Engineer)

## License

---

---

This project is licensed under the MIT License - see the LICENSE ↗ file for details.

## Acknowledgments

---

---

- Kaggle Housing Dataset
- Scikit-learn and XGBoost communities
- Optuna optimization framework
- MLflow experiment tracking
- Evidently ML model monitoring

## Contact

---

---

For questions or support, please open an issue on GitHub or contact ↗ the maintainers.

---

... note::

(continues on next page)

(continued from previous page)

This documentation is comprehensive and up-to-date as of December ~~2025~~. For the latest updates, please check the GitHub repository.

Last Updated: December 6, 2025