

# VNIT Housing Price Prediction Documentation

**version**

**2025, Shivam Awasthi & Abhishek Bhaduria**

December 07, 2025



# Contents

VNIT Housing Price Prediction	1
Project Overview	1
Key Metrics	1
Quick Start	1



# VNIT Housing Price Prediction

Python 3.8+

License MIT

Welcome to the VNIT Housing Price Prediction project documentation. This comprehensive guide covers the complete machine learning pipeline including data preprocessing, exploratory data analysis, model optimization, and deployment.

## Project Overview

The VNIT Housing Price Prediction project is an end-to-end machine learning solution that predicts house prices using advanced regression techniques. The project implements best practices in data science including:

- **Data Preprocessing & EDA** - Comprehensive data cleaning and exploration
- **Hyperparameter Optimization** - Optuna-based parameter tuning
- **Model Training** - XGBoost with cross-validation
- **Experiment Tracking** - MLflow for reproducibility
- **Performance Monitoring** - Evidently for model health checks
- **Workflow Orchestration** - Prefect for pipeline management

## Key Metrics

- **Model R<sup>2</sup> Score:** 0.8990
- **RMSE:** \$27,836.04
- **MAE:** \$17,279.64
- **Optimization Trials:** 50 (Optuna)
- **Cross-Validation Folds:** 5

## Quick Start

### 1. Install Dependencies:

```
pip install -r requirements.txt
```

### 2. Run Preprocessing Pipeline:

```
=====
VNIT Housing Price Prediction
=====
```

```
.. image:: https://img.shields.io/badge/Python-3.8+-blue.svg
.. image:: https://img.shields.io/badge/License-MIT-green.svg
```

Welcome to the VNIT Housing Price Prediction project documentation. This comprehensive g

Project Overview

```
=====
```

The VNIT Housing Price Prediction project is an end-to-end machine learning solution tha

- **\*\*Data Preprocessing & EDA\*\*** - Comprehensive data cleaning and exploration
- **\*\*Hyperparameter Optimization\*\*** - Optuna-based parameter tuning

## VNIT Housing Price Prediction

- **Model Training** - XGBoost with cross-validation
- **Experiment Tracking** - MLflow for reproducibility
- **Performance Monitoring** - Evidently for model health checks
- **Workflow Orchestration** - Prefect for pipeline management

### Key Metrics

=====

- **Model R<sup>2</sup> Score**: 0.8990
- **RMSE**: \$27,836.04
- **MAE**: \$17,279.64
- **Optimization Trials**: 50 (Optuna)
- **Cross-Validation Folds**: 5

### Quick Start

=====

#### 1. **Install Dependencies**:

```
pip install -r requirements.txt
```

#### 2. **Run Preprocessing Pipeline**:

```
cd src/preprocessing
python preprocessing.py ../../data/train.csv
```

#### 3. **Train Model**:

```
cd src/modeling
python modeling.py
```

#### 4. **Launch Streamlit App**:

```
streamlit run src/app/streamlit_app.py
```

### Project Structure

=====

::

```
VNIT_project/
    data/
        train.csv          # Dataset files
        test.csv           # Training data
        X_train.csv        # Test data
        data_description.txt # Preprocessed features
        models/
            xgboost_model.joblib # Feature descriptions
        src/
            preprocessing/
                preprocessing.py   # Trained models
                monitoring.py     # Serialized XGBoost model
                EDADataPreProcessing.ipynb
            modeling/
                modeling.py       # Data preprocessing modules
                Modeling.ipynb    # Preprocessing pipeline
            app/
                streamlit_app.py  # Data quality monitoring
            mlruns/             # Model training modules
                                # Model pipeline
                                # Streamlit application
                                # MLflow experiment tracking
```

```
■■■ docs/                                # Documentation
■■■ Dockerfile                            # Docker configuration
■■■ docker-compose.yml                     # Docker Compose setup
■■■ requirements.txt                      # Python dependencies
```

## Table of Contents

```
=====
```

```
.. toctree::
:maxdepth: 2
:caption: Getting Started
```

```
getting_started
installation
quick_start
```

```
.. toctree::
:maxdepth: 2
:caption: Modules & APIs
```

```
modules/preprocessing
modules/monitoring
modules/modeling
```

```
.. toctree::
:maxdepth: 2
:caption: Guides
```

```
guides/data_preprocessing
guides/model_training
guides/deployment
```

```
.. toctree::
:maxdepth: 2
:caption: Reference
```

```
api_reference
troubleshooting
faq
```

## Features

```
=====
```

- ✓ \*\*Automated Data Pipeline\*\* - End-to-end preprocessing with Prefect
- ✓ \*\*Hyperparameter Optimization\*\* - Optuna integration for parameter tuning
- ✓ \*\*Experiment Tracking\*\* - MLflow for model versioning and comparison
- ✓ \*\*Data Quality Monitoring\*\* - Evidently for drift detection
- ✓ \*\*Model Evaluation\*\* - Comprehensive metrics and validation
- ✓ \*\*Production Ready\*\* - Docker support and API deployment
- ✓ \*\*Documentation\*\* - Complete Sphinx documentation with examples

## Technologies Used

```
=====
```

- \*\*ML Framework\*\*: XGBoost
- \*\*Hyperparameter Tuning\*\*: Optuna
- \*\*Experiment Tracking\*\*: MLflow
- \*\*Workflow Orchestration\*\*: Prefect
- \*\*Data Quality\*\*: Evidently

# VNIT Housing Price Prediction

- **Data Processing**: Pandas, NumPy, Scikit-learn
- **Visualization**: Matplotlib, Seaborn
- **Deployment**: Streamlit, Docker

## Performance

=====

The final XGBoost model achieves excellent performance on the test set:

- **R<sup>2</sup> Score**: 0.8990 (explains 89.90% of variance)
- **Root Mean Squared Error**: \$27,836
- **Mean Absolute Error**: \$17,280
- **Training Samples**: 1,168
- **Test Samples**: 292

## Contributing

=====

Contributions are welcome! Please feel free to submit pull requests or open issues for b

## Authors

=====

- **Shivam Awasthi** (Data Science Engineer)
- **Abhishek Bhadauria** (Data Science Engineer)

## License

=====

This project is licensed under the MIT License - see the LICENSE file for details.

## Acknowledgments

=====

- Kaggle Housing Dataset
- Scikit-learn and XGBoost communities
- Optuna optimization framework
- MLflow experiment tracking
- Evidently ML model monitoring

## Contact

=====

For questions or support, please open an issue on GitHub or contact the maintainers.

---

.. note::

This documentation is comprehensive and up-to-date as of December 2025. For the latest

Last Updated: December 6, 2025